

C++ Exercise WS15/16



UNIVERSITÄT
HEIDELBERG
ZUKUNFT
SEIT 1386

Project: NUMA Migration Martin Wenzel, Dennis Sebastian Rieber

Structure



UNIVERSITÄT
HEIDELBERG
ZUKUNFT
SEIT 1386


- Introduction
 - What is NUMA?
 - Motivation – Why are we doing this?
 - State of the Art – What else is out there?
- Realization – What can we do about the Problem?
 - Goals and Concepts
 - Interface
 - Implementation
- Verification – Was it all worth it?
 - Migration
 - Stencil
- Summary

What is NUMA?



UNIVERSITÄT
HEIDELBERG
ZUKUNFT
SEIT 1386

- Non-Uniform Memory Architecture
- Multisocket Systems
- shared memory systems
- cache coherent



P Prozessor
M Speicher
VN Verbindungsnetzwerk


NUMA-Parallelrechner

Motivation – Why are doing this?



UNIVERSITÄT
HEIDELBERG
ZUKUNFT
SEIT 1386

- Remote data access
- first touch policy
 - Employed by all modern operating systems
 - Initializes memory page on the domain of the thread, which is the first to touch the data



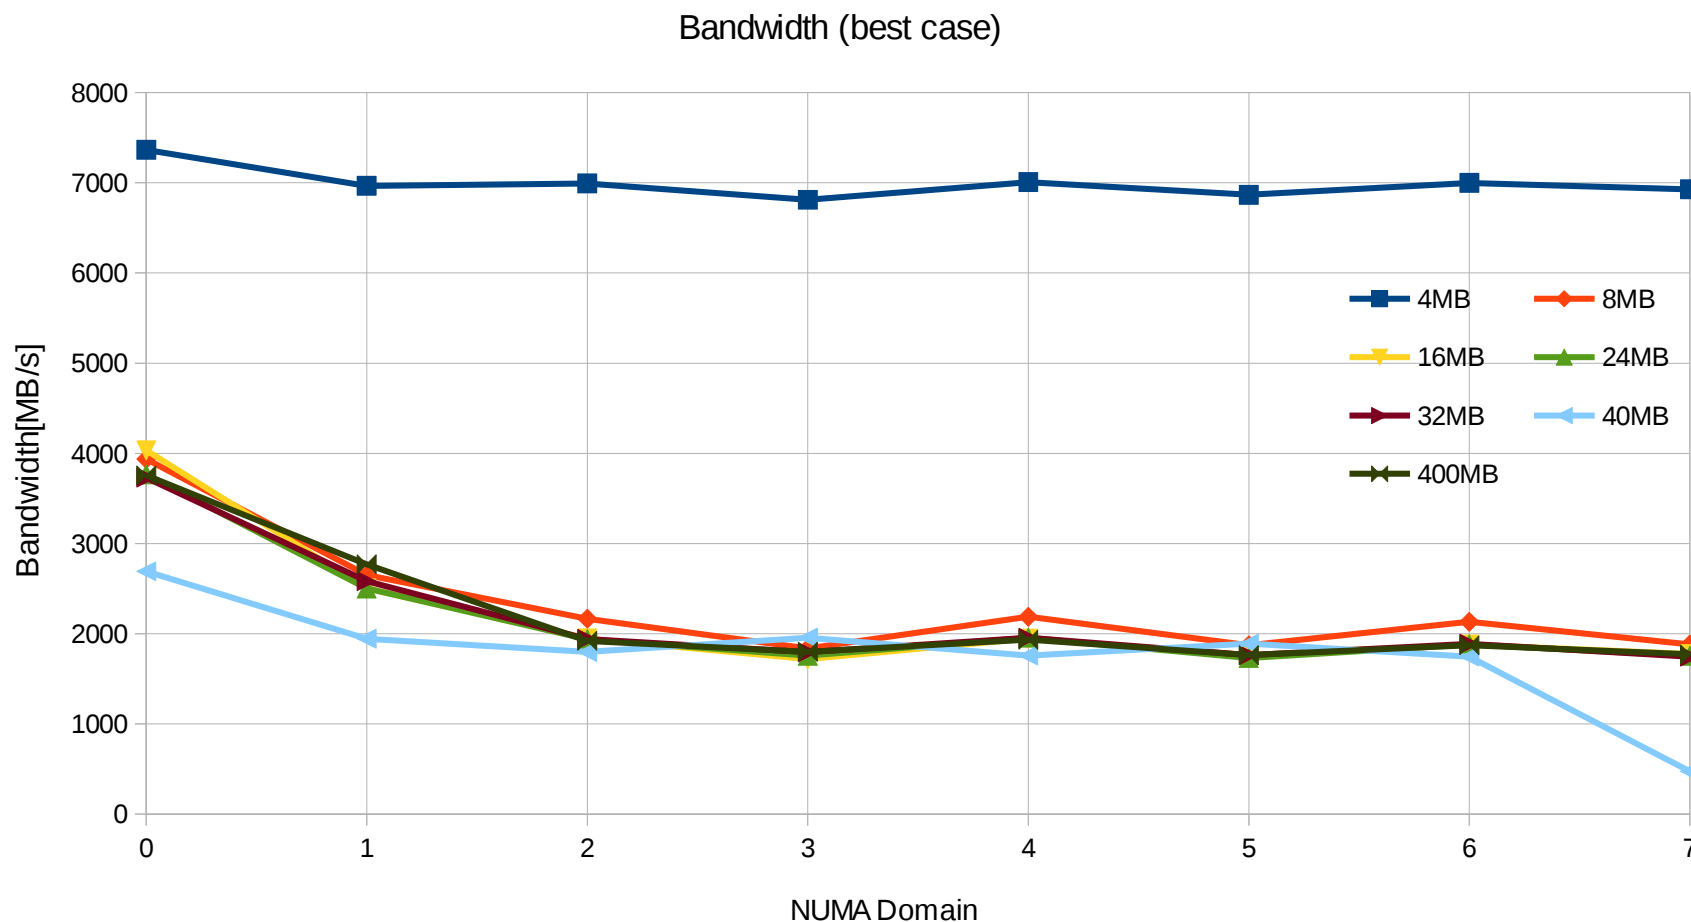
P Prozessor
M Speicher
VN Verbindungsnetzwerk

NUMA-Parallelrechner

Motivation – Costs of remote memory access



UNIVERSITÄT
HEIDELBERG
ZUKUNFT
SEIT 1386



State of the Art – What else is out there?



UNIVERSITÄT
HEIDELBERG
ZUKUNFT
SEIT 1386

- Hwloc – hardware locality
 - Explores system architecture
 - Creates architectural model of the system
 - Huge interface
- Libnuma
 - Only available on Unix systems
 - Classic C style interface
 - Memory page and thread migration possibilities

State of the Art – libnuma



UNIVERSITÄT
HEIDELBERG
ZUKUNFT
SEIT 1386

```
#include <numa.h>
int      _pagesize          = sysconf(_SC_PAGESIZE);
int      _numa_pagecount    = _size * sizeof(float) / _pagesize;
int*     _numa_destination   = new int[_numa_pagecount];
int*     _status            = new int[_numa_pagecount];
void**   _numa_pages         = (void**) malloc (_numa_pagecount
*sizeof(void *));

for (int i = 0; i < _numa_pagecount; i++){
    _numa_pages[i] = &_data[i*( _pagesize / sizeof(float))];
    _numa_destination[i] = domain;
}

if (move_pages(0, _numa_pagecount, _numa_pages,
_numa_destination, _status, MPOL_MF_MOVE) != 0) {
    std::cout << "something went wrong" << std::endl;
}
```

Structure



UNIVERSITÄT
HEIDELBERG
ZUKUNFT
SEIT 1386

- Introduction
 - What is NUMA?
 - Motivation – Why are we doing this?
 - State of the Art – What else is out there?
- Realization – What can we do about the Problem?
 - Goals and Concepts
 - Interface
 - Implementation
- Verification – Was it all worth it?
 - Migration
 - Stencil
- Summary

Goals



UNIVERSITÄT
HEIDELBERG
ZUKUNFT
SEIT 1386

- Core
 - ✓ Container
 - ✓ First touch migration of complete data block
- Target
 - ✓ Distribute pages over several domains
 - ✓ Templated Container
- Enhanced
 - ✗ Intelligent distribution based on access patterns
 - ✗ Complete portability

Concepts



UNIVERSITÄT
HEIDELBERG
ZUKUNFT
SEIT 1386

- NUMA aware container class
 - Knowing the available domains
 - Move data between domains
- STL-like behavior
 - Iterators
 - Support foreach Syntax
 - Templates
- Using first-touch to move data, no syscalls required
- easy user Interface with c++ language features

Realization - Interface



```
template <class T> class numavec
```

```
    typedef typename std::vector<T>::iterator iterator;  
    typedef typename std::vector<T>::const_iterator const_iterator  
    typedef std::initializer_list<int> init;
```

```
    numavec<T>( int );  
    numavec<T>(int, std::map<int, int> &);  
    ~numavec<T>() = default;
```

```
    numavec<T>( const numavec& ); // copy  
    numavec<T>( numavec&& ); // move  
    numavec<T>& operator=( const numavec& ); // copy  
    numavec<T>& operator=( numavec&& ); // move
```

```
    T& operator[](int in) { return _data[in]; };  
    T& at(int);
```

```
    int size() const { return _data.size(); };  
    const T* data() const { return _data; };
```

```
    void migrate_to(int);  
    void distribute( init );  
    void uneven_distribute( std::map<int,int> &);
```

```
    const_iterator cbegin() const { return _data.cbegin(); };  
    const_iterator cend() const { return _data.cend() + _size; };
```

```
    const_iterator begin() const {return _data.begin();};  
    const_iterator end() const { return _data.end() + _size; };
```

```
    iterator begin() {return _data.begin();};  
    iterator end() { return _data.end() + _size;};
```

Realization



UNIVERSITÄT
HEIDELBERG
ZUKUNFT
SEIT 1386

```
_core_pin( _domain_core[_destination.begin()->second]);

std::vector<T> new_data;

new_data.swap(_data);

_data.clear();
_data.shrink_to_fit();
_data.reserve(_size);

std::vector<std::thread> threads(_destination.size());
    int pos1 = 0; int i = 0;
    for (auto& e: _destination){ //-- iterate each interval
        threads[i] = std::thread(&numavec::_migrate_part,this, e.second, pos1,
                                e.first, std::ref(new_data));
        pos1 = e.first + 1; //-- define the next start
        ++i;
    }
for ( unsigned int i = 0; i < threads.size(); i++ ) {
    threads[i].join();
}
```

Realization – Thread pin



UNIVERSITÄT
HEIDELBERG
ZUKUNFT
SEIT 1386

```
void numavec<T>::_core_pin( int core ) {  
    cpu_set_t cpuset;  
    CPU_ZERO(&cpuset);  
    CPU_SET( core , &cpuset);  
    pthread_setaffinity_np(pthread_self(), sizeof(cpu_set_t),  
&cpuset);  
}
```

Structure



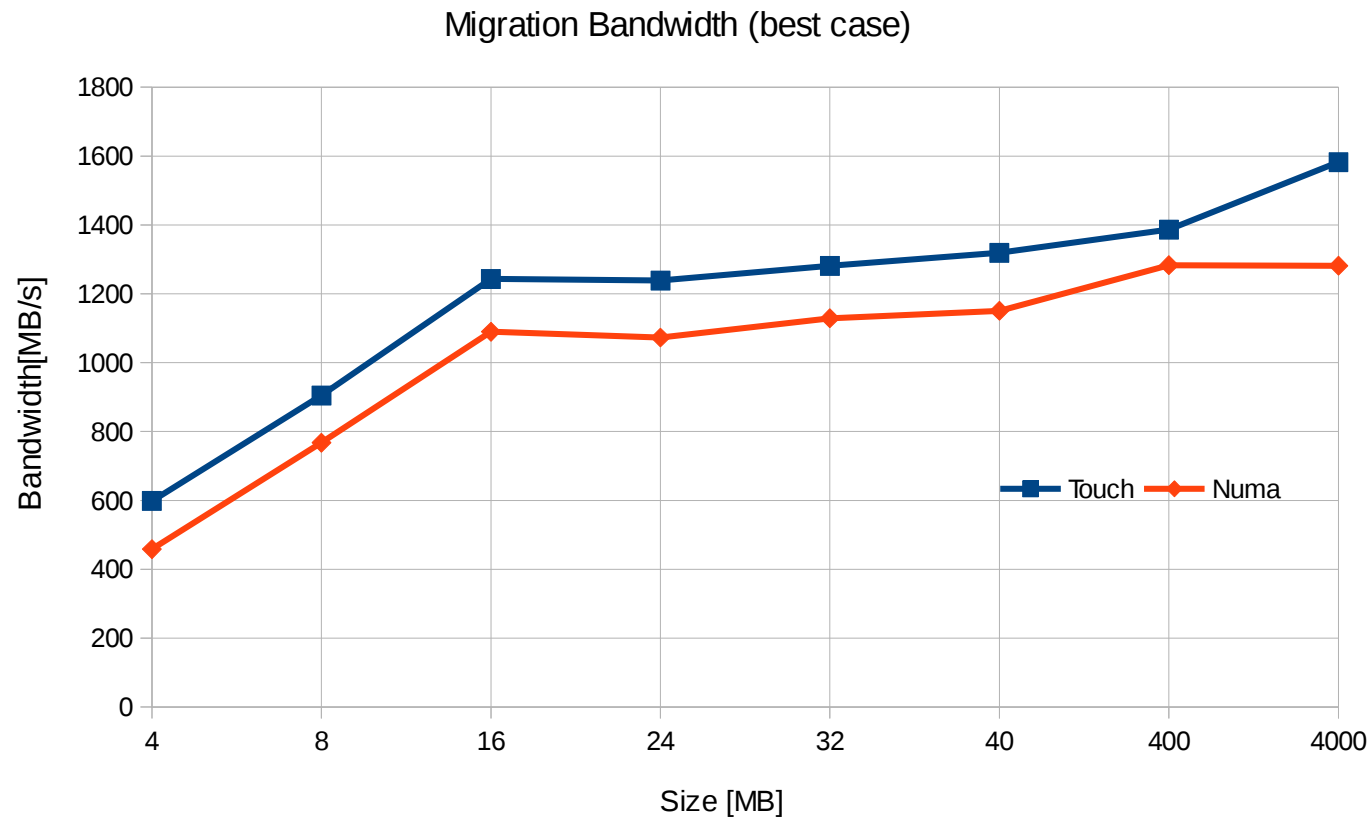
UNIVERSITÄT
HEIDELBERG
ZUKUNFT
SEIT 1386

- Introduction
 - What is NUMA?
 - Motivation – Why are we doing this?
 - State of the Art – What else is out there?
- Realization – What can we do about the Problem?
 - Goals and Concepts
 - Interface
 - Implementation
- Verification – Was it all worth it?
 - Migration
 - Stencil
- Summary

Verification - Migration



UNIVERSITÄT
HEIDELBERG
ZUKUNFT
SEIT 1386



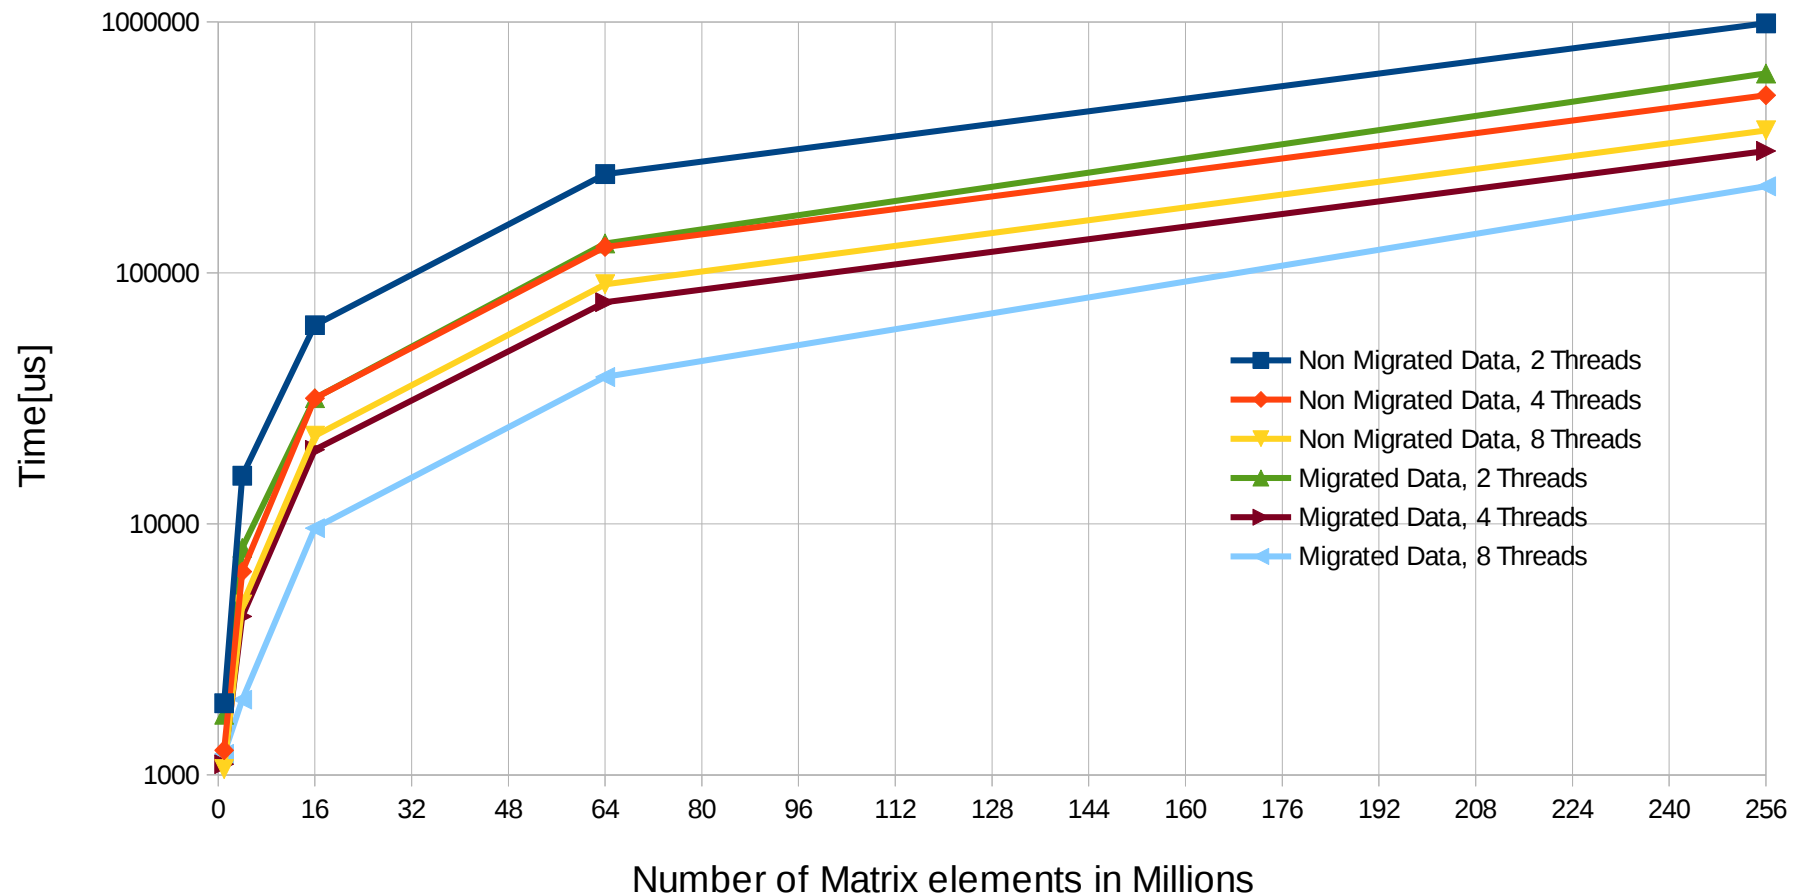
Verification – Stencil



UNIVERSITÄT
HEIDELBERG
ZUKUNFT
SEIT 1386

Stencil Benchmark

Avg. Time/Iteration



Structure



UNIVERSITÄT
HEIDELBERG
ZUKUNFT
SEIT 1386

- Introduction
 - What is NUMA?
 - Motivation – Why are we doing this?
 - State of the Art – What else is out there?
- Realization – What can we do about the Problem?
 - Goals and Concepts
 - Interface
 - Implementation
- Verification – Was it all worth it?
 - Migration
 - Stencil
- Summary

Summary



UNIVERSITÄT
HEIDELBERG
ZUKUNFT
SEIT 1386

- Achieved core and target goals
- Created easy to use container that hides NUMA complexity
- Future Goals
 - Intelligent distribution
 - Portable Version
 - Separate NUMA administration and Container administration

C++11 Features



UNIVERSITÄT
HEIDELBERG
ZUKUNFT
SEIT 1386

- Mutex
- Move Semantics
- Threads
- Initializer lists
- Constructor delegation